

Ex.No: 8	
Date: 15.9.25	

**Implementing a Combiner in Word Count****Aim:**

To implement a Combiner in Hadoop MapReduce for Word Count, and analyze its effect on performance and intermediate data transfer.

**Procedure:****1. Set up environment**

- Ensure Hadoop (3.x) is installed and HDFS/YARN is running in pseudo-distributed or cluster mode.
- Java (JDK 1.8+) and Hadoop client libraries should be on your classpath.

**2. Create three Java source files**

- WC\_Mapper.java — Mapper class
- WC\_Reducer.java — Reducer and Combiner
- WC\_Driver.java — Job configuration and submission

**3. Build the JAR**

- Use Maven, Ant, or javac + jar to package the classes into wordcount-combiner.jar.

**4. Prepare input**

- On your local machine, create input text (e.g., input.txt).
- Upload to HDFS:
 

```
hadoop fs -mkdir /wordcount/input
hadoop fs -put input.txt /wordcount/input
```

**5. Run the job**

- Command:

```
hadoop jar wordcount-combiner.jar WC_Driver /wordcount/input /wordcount/output
```

**6. Observe output**

- Command

```
hadoop fs -cat /wordcount/output/part-r-00000
```

**7. Analyze results**

- Compare job runtime **with and without** the combiner; measure intermediate map outputs to see how the combiner reduces shuffle traffic.

**Program:**

**WC\_Mapper.java**

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
public class WC_Mapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    @Override
    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

**WC\_Reducer.java**

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class WC_Reducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();
    @Override

```

```
protected void reduce(Text key, Iterable<IntWritable> values, Context context)
    throws IOException, InterruptedException {
    int sum = 0;
    for (IntWritable val : values) {
        sum += val.get();
    }
    result.set(sum);
    context.write(key, result);
}
```

### WC\_Driver.java

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class WC_Driver {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: WC_Driver <input> <output>");
            System.exit(2);
        }
        Job job = Job.getInstance(new Configuration(), "Word Count with Combiner");
        job.setJarByClass(WC_Driver.class);
        job.setMapperClass(WC_Mapper.class);
        job.setCombinerClass(WC_Reducer.class); // Combiner enabled
        job.setReducerClass(WC_Reducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
    }
}
```

```
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

**Output:**

input.txt : hello world hello Hadoop Combiner

**Final reducer output:**

Combiner 1  
Hadoop 1  
hello 3  
world 1

**Result:**

Thus, the program has been executed successfully and the output is also verified.